

Lab 10 Classification: K-NN

Lusine Zilfimian

April 22 (Wednesday), 2020

Contents

- Libraries
- Data Preparation / Understanding the data
- k-NN Classification
- k-NN Regression
- Questions

Needed packages

```
library(caret) # for createDataPartition() and trainControl(), train()
library(class) # for knn()
library(dplyr) # for data manipulation
library(MASS) # for Boston Dataset
library(FNN) # for knn.reg()
library(ggvoronoi) # for stat_voronoi()
library(ggplot2) # ok, you know it
```

Data Preparation / Understanding the data

- Read the data Diabetes.csv

```
diab <- read.csv("Diabetes.csv")
str(diab)
```

```
## 'data.frame':    768 obs. of  9 variables:
## $ NTS : int  6 1 8 1 0 5 3 10 2 8 ...
## $ PGC : int  148 85 183 89 137 116 78 115 197 125 ...
## $ DBP : int  72 66 64 66 40 74 50 0 70 96 ...
## $ TSFT : int  35 29 0 23 35 0 32 0 45 0 ...
## $ INS : int  0 0 0 94 168 0 88 0 543 0 ...
## $ BMI : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
## $ DPF : num  0.627 0.351 0.672 0.167 2.288 ...
## $ Age : int  50 31 32 21 33 30 26 29 53 54 ...
## $ Class: Factor w/ 2 levels "Negative","Positive": 2 1 2 1 2 1 2 1 2 2
```


Understanding the data

Source: <https://www.kaggle.com/uciml/pima-indians-diabetes-database>

- NTS - Number of times pregnant
- PGC - Plasma glucose concentration a 2 hours in an oral glucose tolerance test
- DBP - Diastolic blood pressure (mm Hg)
- TSFT - Triceps skin fold thickness (mm)
- INS - Hour serum insulin (μ U/ml)
- BMI - Body mass index ($\text{weight in kg}/(\text{height in m})^2$)
- DPF - Diabetes pedigree function
- Age (years)
- Class variable

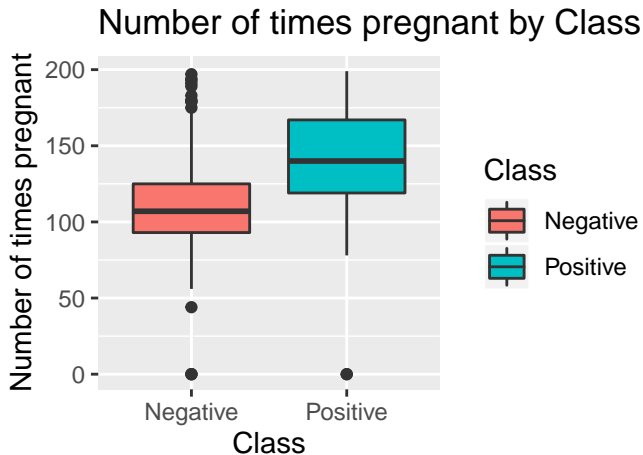
Understanding the data

```
ggplot(data = diab, aes(x = Class, y = DBP, fill = Class)) +  
  geom_boxplot() + labs(x = "Class", y = "Blood pressure",  
    title = "Blood pressure by Class")
```



Understanding the data

```
ggplot(data = diab, aes(x = Class, y = PGC, fill = Class))+  
  geom_boxplot() + labs(x = "Class", y = "Number of times pregnant",  
    title = "Number of times pregnant by Class")
```



Understanding the data

```
ggplot(data = diab, aes(x = Class, y = BMI, fill = Class))+  
  geom_boxplot() + labs(x = "Class", y = "Body mass index",  
    title = "Body mass index by Class")
```



Data Preparation

- Divide the dataset into training and testing sets:

```
set.seed(2708)
split <- diab$Class %>% createDataPartition(p = 0.8, list = FALSE)

train.data <- diab[split, ]
test.data <- diab[-split, ]
```

k-NN Classification (Sample Validation)

- Make prediction on the testing set (single case):
- The function arguments:
 - 1 train: The training set without actual label
 - 2 test: the testing set without actual label
 - 3 cl: actual class labels from Train set
 - 4 k: the number of neighbors to consider

```
# Make prediction on the testing set (one observation)  
(knn0 <- knn(train = train.data[,-9], test = test.data[1,-9],  
  cl = train.data$Class, k = 5))
```

```
## [1] Negative  
## attr("nn.index")  
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]  604  355  382  126  606  
## attr("nn.dist")  
##      [,1]      [,2]      [,3]      [,4]      [,5]  
## [1,] 10.55812 19.19861 22.04193 23.60764 24.60926  
## Levels: Negative
```

k-NN Classification (Sample Validation)

- Make prediction for the whole Test dataset.
- The first column of each dataframe is the actual class, so we take it out

```
knn1 <- knn(train = train.data[, -9], test = test.data[, -9],  
  cl = train.data$Class, k=5)  
length(knn1) == dim(test.data)[2]
```

```
## [1] FALSE
```

```
head(knn1, 10)
```

```
## [1] Negative Negative Positive Negative Positive Negative Negative  
## [8] Negative Negative Negative  
## Levels: Negative Positive
```

k-NN Classification (Sample Validation)

- Lets look at the accuracy

```
table(knn1, test.data$Class)
```

```
##
```

```
## knn1      Negative Positive
```

```
##   Negative      83      29
```

```
##   Positive      17      24
```

```
# Accuracy
```

```
(83 + 24)/153
```

```
## [1] 0.6993464
```

```
mean(knn1 == test.data$Class)
```

```
## [1] 0.6993464
```

```
table(test.data$Class)/153
```

```
##
```

```
##   Negative Positive
```

```
## 0.6535948 0.3464052
```

k-NN Classification (Sample Validation)

- You can use knn function to predict probabilities of belonging to a class as well.
- Use argument prob = T in the function
- Access the probabilities using function attr, inside giving the object and attribute name "prob"
- Case 1: there were 5 neighbors, 4 of them were Negative cases so we predict the case as being Negative with the probability of 0.8 (4/5)

```
knn2 <- knn(train = train.data[, -9], test = test.data[, -9],  
            cl = train.data$Class, k = 6, prob = T)  
df <- data.frame(class = knn1, probs=attr(knn2, "prob"))  
head(df)
```

```
##      class      probs  
## 1 Negative 0.6666667  
## 2 Negative 0.8333333  
## 3 Positive 0.6666667  
## 4 Negative 0.6666667  
## 5 Positive 0.6666667  
## 6 Negative 0.8333333
```

k-NN Classification (Cross Validation)

- `trainControl` creates object that will be used in modeling. In this case we specify that we want to do cross-validation (`method="cv"`) with 10 folds (`number=10`).

```
set.seed(2708)
ctrl <- trainControl(method="cv", number = 10)
```

Function `train` from `caret`

- give the formula
- give the data
- set the method for training (`method="knn"`)
- `trControl` specifies the training control object
- `preProcess`: calculate the Z-scores
- `tuneLength`: for how many parameters of `k` do you want to build model

k-NN Classification (Cross Validation)

The `tuneLength` parameter tells the algorithm to try different **default** values for the main parameter

```
set.seed(2708)
knn_c <- train( Class~., data = diab, method = "knn",
  trControl = ctrl, preProcess = c("center","scale"), tuneLength = 10)
```

k-NN Classification (Cross Validation)


- The choice of parameters is always the same. The column k shows the number of neighbours that are used during the training
- Column Accuracy shows the average accuracy on the testing sets (total 10 testing sets)
- Overall 10 different values for k are used for modeling

```
knn_c$results[,1:2]
```

##	k	Accuracy
## 1	5	0.7422078
## 2	7	0.7409262
## 3	9	0.7369446
## 4	11	0.7382946
## 5	13	0.7331340
## 6	15	0.7384142
## 7	17	0.7475051
## 8	19	0.7526658
## 9	21	0.7578776
## 10	23	0.7513500

k-NN Classification (Cross Validation)

```
plot(knn_c)
```




```
knn_c$bestTune
```

k-NN Classification (Cross Validation)

- To specify vectors of k's for the model, use `tuneGrid` (not `tuneLenght`)

```
grid <- expand.grid(k = 10:20); set.seed(2708)
knn_c1<-train(Class~., data = diab, method="knn", trControl = ctrl,
  preProcess = c("center","scale"), tuneGrid = grid)
plot(knn_c1)
```



KNN Regression (Sample Validation)

This dataset contains information collected by the U.S Census Service concerning housing in the area of Boston Mass.


```
set.seed(42)
idx <- sample(1:nrow(Boston), size = 250)
trn_boston <- Boston[idx, ]
tst_boston <- Boston[-idx, ]
```

```
RMSEtrain <- NULL
RMSEtest <- NULL
for (i in 1:100){
  trainknn <- knn.reg(train = trn_boston[, -1],
                      test = trn_boston[, -1], y = trn_boston$crim, k = i)
  testknn <- knn.reg(train = trn_boston[, -1], test = tst_boston[, -1],
                     y = trn_boston$crim, k = i)
  RMSEtrain[i] <- sqrt(mean((trn_boston$crim - trainknn$pred)^2))
  RMSEtest[i] <- sqrt(mean((tst_boston$crim - testknn$pred)^2))
}
```

KNN Regression (Sample Validation)


```
df <- data.frame(x = 1:100, train = RMSEtrain, test = RMSEtest)
ggplot(data = df, aes(x = x, y = train)) +
  geom_point(color = "blue")+
  labs(y = "Train RMSE", x = "k", title="Train RMSE")
```

Train RMSE



KNN Regression (Cross Validation)

```
grid <- expand.grid(k = 1:20); set.seed(1)
model <- train( crim~., data = Boston, method = "knn",
  trControl = trainControl("cv", number = 10),
  preProcess = c("center", "scale"), tuneGrid = grid)
plot(model)
```



KNN Regression (Cross Validation)

```
head(model$results[,c(1,2)], n=10)
```


```
##      k      RMSE
## 1    1  7.070863
## 2    2  6.815805
## 3    3  6.122864
## 4    4  5.811194
## 5    5  5.872013
## 6    6  5.767924
## 7    7  5.686235
## 8    8  5.446108
## 9    9  5.438080
## 10  10  5.451263
```

```
model$results[which.min(model$results$RMSE),]
```

```
##      k      RMSE  Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 12  12  5.436749  0.6297794  1.804765  3.145291  0.2207016  0.6775957
```

Voronoi Tessellation


```
set.seed(2708)
x <- sample(1:200,100); y <- sample(1:200,100)
z <- rbinom(n = 100, size = 1, prob = 0.6); df <- data.frame(x,y,z)
ggplot(df, aes(x,y, col=factor(z))) + stat_voronoi(geom="path") +
  geom_point() + theme(legend.position = "none")
```



Question

- ❶ If for B the nearest is A, does it mean that for A the nearest is B

```
x1 <- c(3,4,4,5,5,6,6,7,15,25,27); x2 <- c(5, 3, 7, 1,9,3,7,5, 5, 5,7 )  
y <-c(rep("o", 7),"A", "B", rep("-", 2)); df <- data.frame(x1, x2, y)  
ggplot(df, aes ( x = x1, y = x2, col = factor(y))) + geom_point()
```



Question

② Problem: What to do?

```
x<-c(4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4)
y<-c(5, 3, 7, 1, 2, 9, 8, 6, 4, 0.5, -5)
df <- data.frame(x,y);
ggplot(df, aes ( x= x, y = y)) + geom_point()
```

